

Discrete-Event Simulation using R-Simmer

KuVS Fachgespräch – WueWoWAS'22 – Würzburg



Stefan Geißler

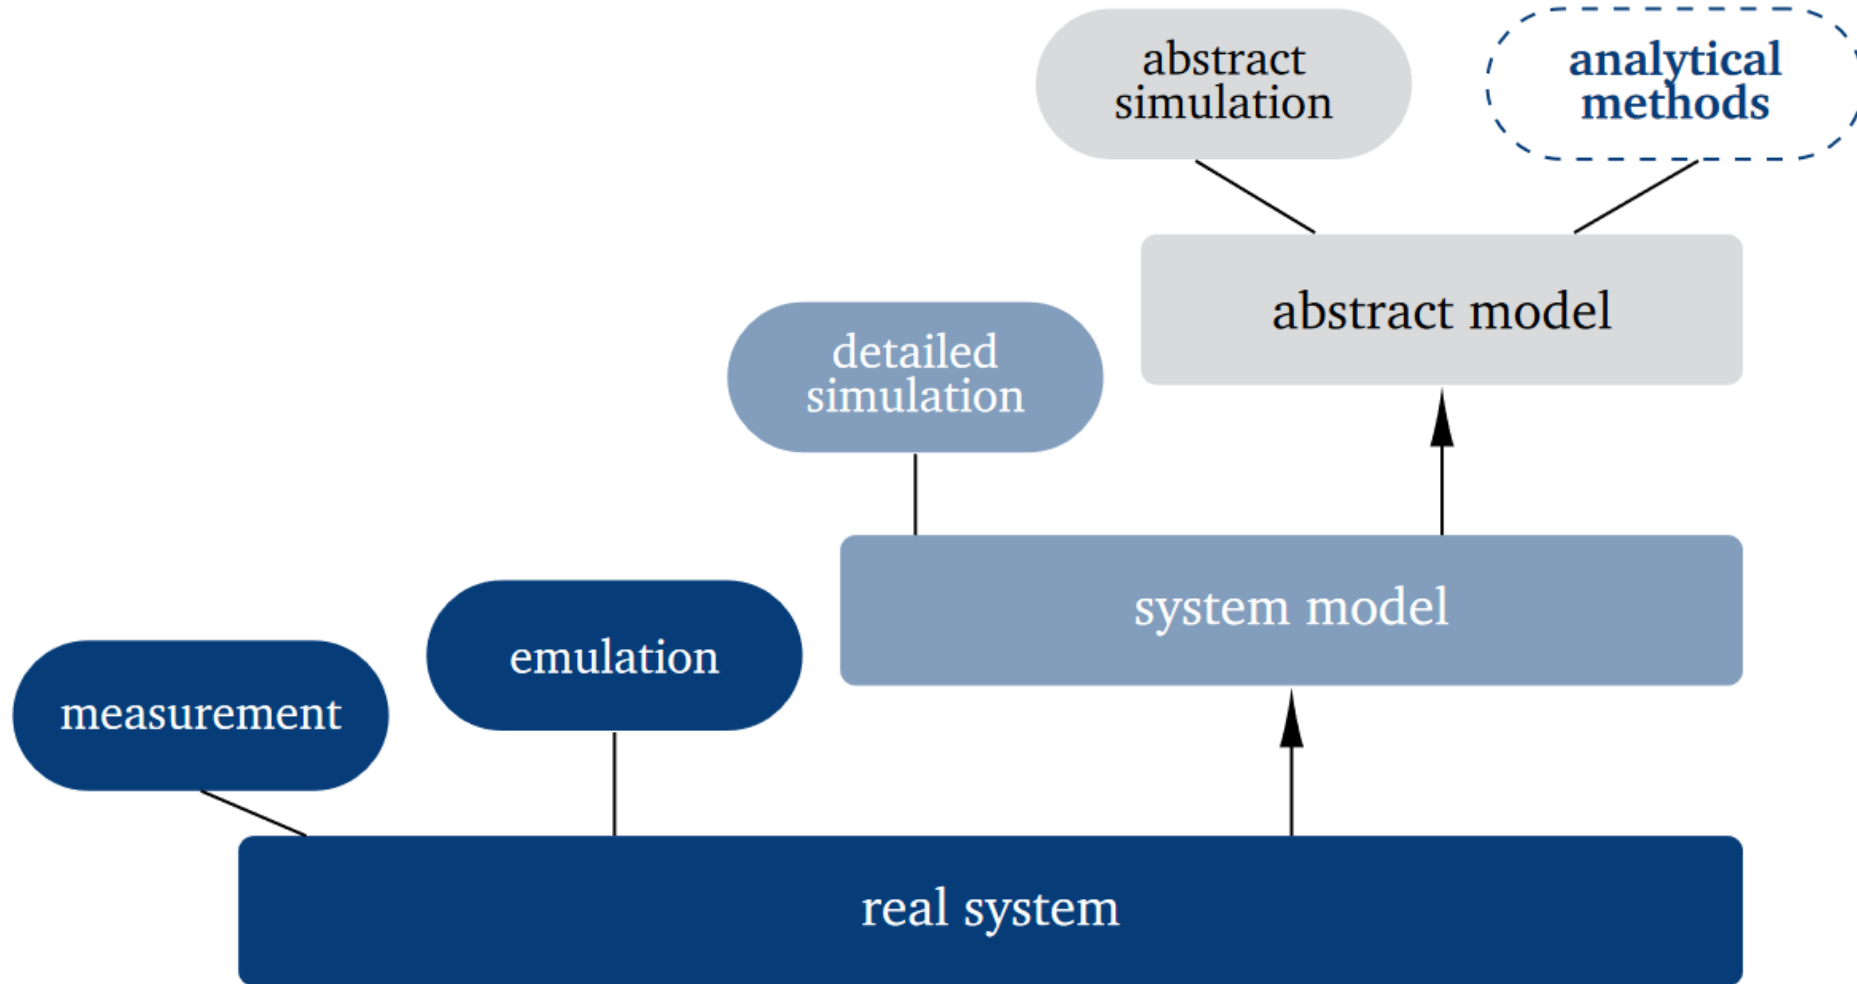
stefan.geissler@uni-wuerzburg.de



From R. Shannon (1975), simulation is

*the process of designing a **model of a real system** and conducting experiments with this model for the purpose either of **understanding the behavior of the system** or of **evaluating various strategies** [...] for the operation of the system.*

*Shannon, Robert E. Systems simulation; the art and science.
No. 04; T57. 62, S4.. 1975.*




Tran-Gia, Phuoc, and Tobias Hoßfeld. *Performance Modeling and Analysis of Communication Networks: A Lecture Note*. BoD—Books on Demand, 2021.

Classification of Simulation Models

Static

- ▶ Simulation of a system at exactly one point in time (e.g. Monte Carlo Simulation)



Deterministic

- ▶ Simulation of a system without randomization (e.g. differential equations for chemical processes)

Continuous

- ▶ Simulation of a system that exhibits state changes at continuous points in time (e.g. chemical reactions)

Dynamic

- ▶ Simulation of a systems behavior over time (e.g. queueing models)



Stochastic

- ▶ Simulation of a system under the influence of stochastic processes (e.g. queueing models)

Discrete

- ▶ Simulation of a system that exhibits state changes at discrete points in time (e.g. queueing models)

Law, Averill M., W. David Kelton, and W. David Kelton. *Simulation modeling and analysis*. Vol. 3. New York: McGraw-hill, 2007.

Classification of Simulation Models

Static

- ▶ Simulation of a system at exactly one point in time (e.g. Monte Carlo Simulation)



Deterministic

- ▶ Simulation of a system without randomization (e.g. differential equations for chemical processes)

Continuous

- ▶ Simulation of a system that exhibits state changes at continuous points in time (e.g. chemical reactions)

Dynamic

- ▶ Simulation of a systems behavior over time (e.g. queueing models)



Stochastic

- ▶ Simulation of a system under the influence of stochastic processes (e.g. queueing models)

Discrete

- ▶ Simulation of a system that exhibits state changes at discrete points in time (e.g. queueing models)

Law, Averill M., W. David Kelton, and W. David Kelton. *Simulation modeling and analysis*. Vol. 3. New York: McGraw-hill, 2007.

Classification of Simulation Models

Static Deterministic Continuous

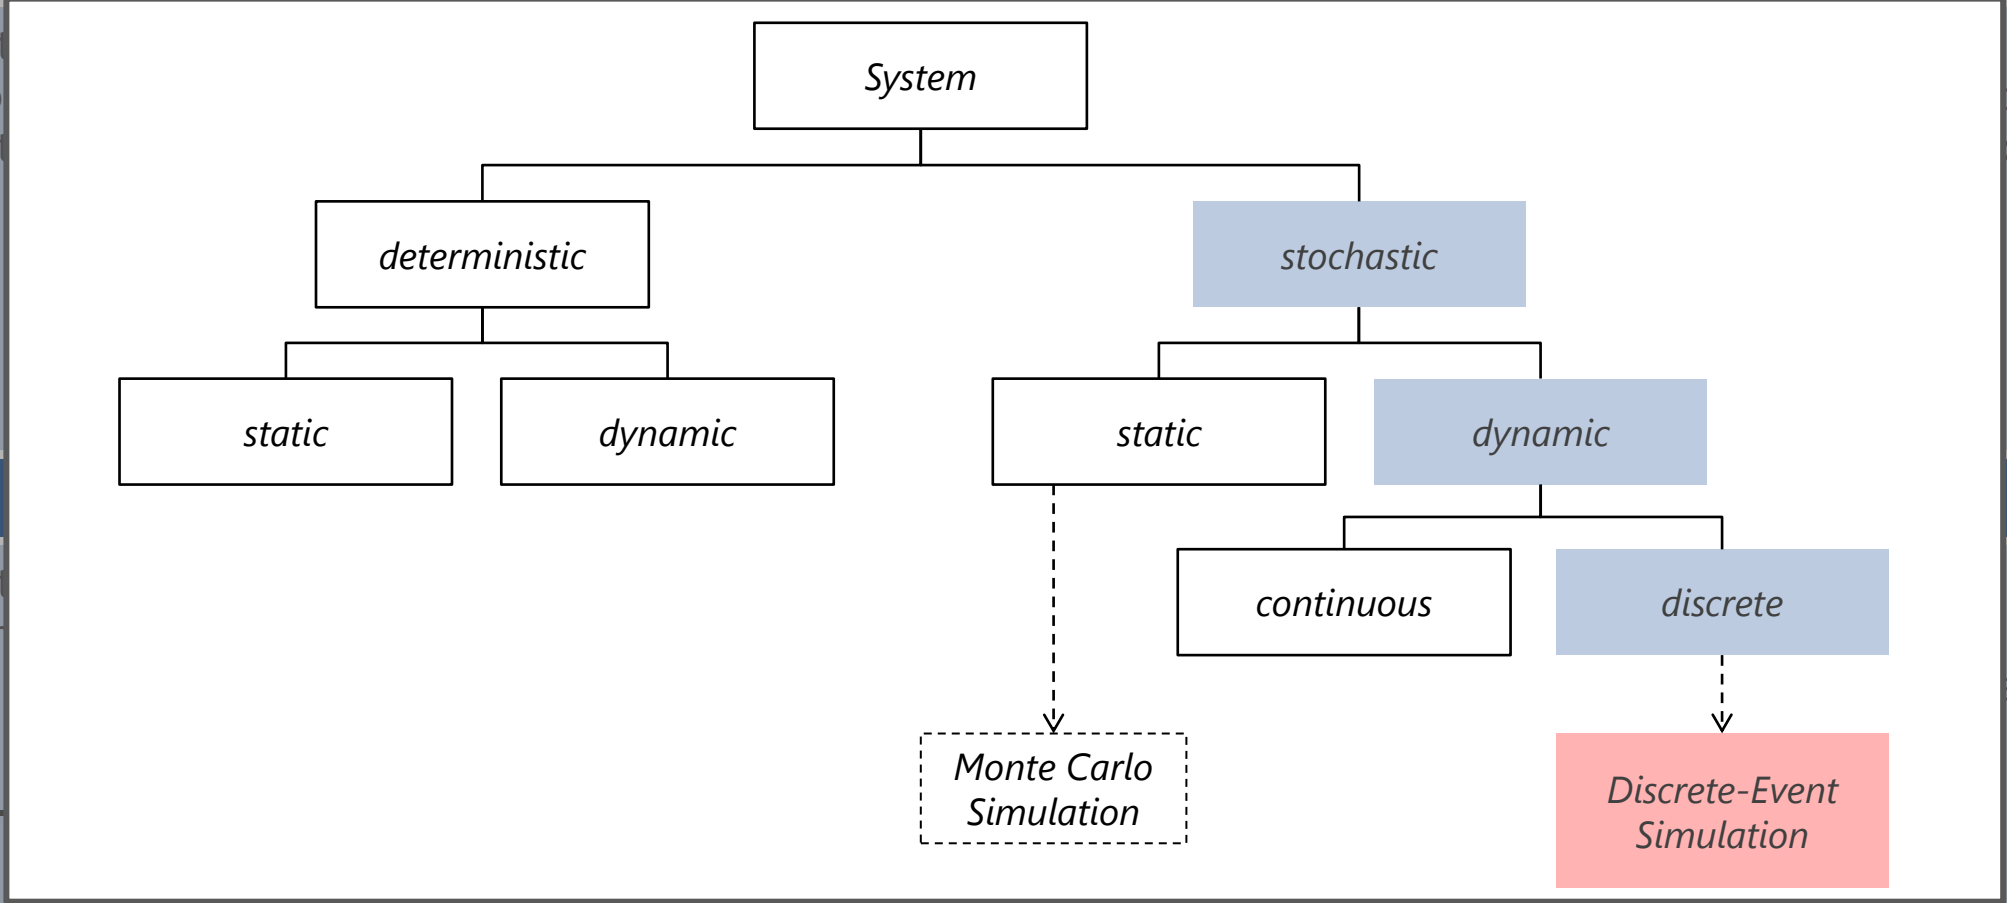
► Simulation of one point in time
Simulation of a system at a single point in time

Dynamic

► Simulation over time
Simulation of a system over time

Simulation of a system that exhibits continuous points in time (e.g., a queueing system)

Simulation of a system that exhibits points in time (e.g., a queueing system)



Law, Averill M., W. David Kelton, and W. David Kelton. Simulation modeling and analysis. Vol. 3. New York: Mcgraw-hill, 2007.

Abstract Idea of a Simulation

- ▶ Base for most simulations is a real world system – e.g. Application, Hardware Component, Distributed System, Industry Process



Random or Periodic Event

- Packet Arrival
- System Signal
- User Interaction

Physical or Logical Component of Real World System

- CPU, Memory
- Network Link
- Human Resource


Near Arbitrary System Reaction Depending on Trigger Type

- Processing End
- Packet Drop
- CPU Interrupt
- Link Congestion
- Event Mishandling

- ▶ Representation of real world system as abstracted sequence of events

→ **(Almost) anything can be simulated!**

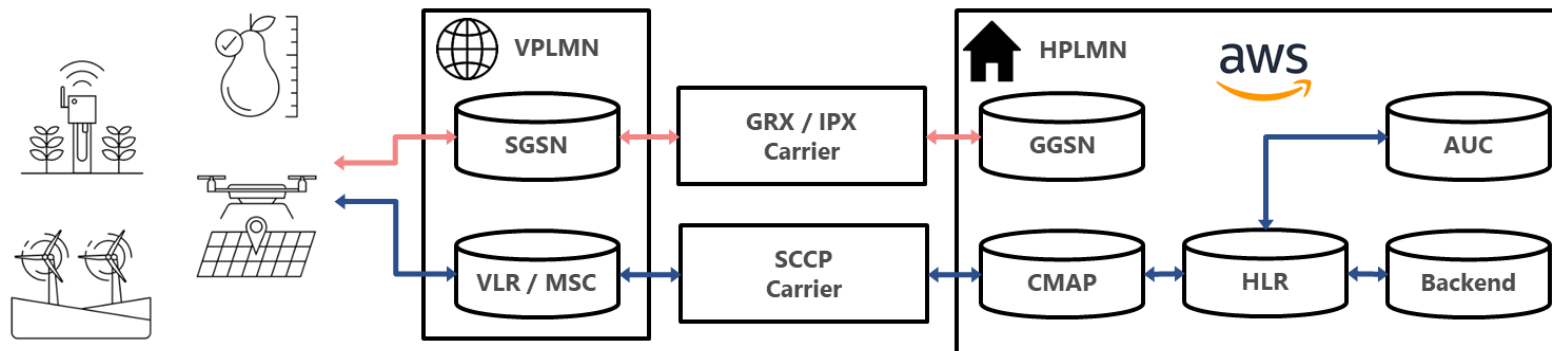
→ **How close to reality can we get?**



Simple Queueing System



Complex Mobile Communication System



simmer: Discrete-Event Simulation for R



Iñaki Ucar

Universidad Carlos III de Madrid

Bart Smeets

dataroots

Arturo Azcorra

Universidad Carlos III de Madrid
IMDEA Networks Institute

Abstract

The **simmer** package brings discrete-event simulation to R. It is designed as a generic yet powerful process-oriented framework. The architecture encloses a robust and fast simulation core written in C++ with automatic monitoring capabilities. It provides a rich and flexible R API that revolves around the concept of *trajectory*, a common path in the simulation model for entities of the same type.

Resources

<https://r-simmer.org>


<https://www.rdocumentation.org/packages/simmer/versions/4.4.2>

<https://groups.google.com/g/simmer-devel>

<https://github.com/r-simmer/simmer>

Advanced concepts (e.g. replication, parallelization):

https://r-simmer.org/slides/20191115_xijur_simmer



- ▶ Events can be triggered at any time
- ▶ Resources have limited **capacity** and may have a limited **queue size**

```
1 library(simmer)
2
3 paket.trajectory <- trajectory() %>%
4   seize("resource") %>%
5   timeout(5) %>%
6   release("resource")
7
8 env <- simmer() %>%
9   add_generator(name = "event", trajectory = paket.trajectory, distribution = at(2,4,6)) %>%
10  add_resource(name = "resource", capacity = 1, queue_size = Inf, mon = T)
11
12 run(env)
13
```

- ▶ **Generators** produce new events
- ▶ **Resources** process events based on their capacity
- ▶ **Trajectories** define interactions between events and resources as well as other events

The simmer API - Trajectories

- ▶ **trajectories** are recipes, or lists, of activities that define the life time of arrivals
- ▶ Similar to **dplyr** for data manipulation

From H. Wickham

[...] by constraining your options, it simplifies how you can think about [something]

▶ Fixed parameters

```
2
3 traj0 <- trajectory() %>%
4   log_("Entering the trajectory") %>%
5   timeout(10) %>%
6   log_("Leaving the trajectory")
7
```

▶ Dynamic parameters

```
8
9 traj1 <- trajectory() %>%
10  log_(function() "Entering the trajectory") %>%
11  timeout(function() 10) %>%
12  log_(function() "Leaving the trajectory")
13
```

```
2
3 paket.trajectory <- trajectory() %>%
4   seize("resource") %>%
5   timeout(5) %>%
6   release("resource")
7
```

Full reference can be found online:
<https://r-simmer.org/reference/>

- ▶ **trajectory()** defines a sequence of actions triggered by an event
 - Is generally processed step by step
 - Can be **rolled back** to create loops
 - Can **branch** or be **cloned**
- ▶ Basic **actions** to use in a trajectory are
 - **seize()** requests capacity from a resource
 - **timeout()** lets time pass
 - **release()** free previously requested capacity
 - **log_()** prints a log message
 - **set_attribute(), get_attribute()** sets/gets an arrival level attribute
 - **set_global(), get_global()** sets/gets a global attribute
 - **rollback()** rolls back the current trajectory by a specific amount of actions


The simmer API – Generators and Resources

```
7  
8 env <- simmer() %>%  
9   add_generator(name = "event", trajectory = paket.trajectory, distribution = at(2,4,6)) %>%  
10  add_resource(name = "resource", capacity = 1, queue_size = Inf, mon = T)  
11
```

- ▶ **add_generator()** defines a source of events or arrivals
 - **name** is the canonical name of the resource
 - **trajectory** defines what generated events do
 - **distribution** dictates when new events occur
- ▶ **add_resource()** defines a claimable resource
 - **name** is the canonical name of the resource
 - **capacity** is the number of available processing
 - **queue_size** is the maximum number of elements allowed to wait before rejects occur
 - **mon** defines whether the resource needs to be monitored

The Gi/Gi/1- ∞ Queue

$$\begin{aligned} a(4) &= 0.4, & a(8) &= 0.5, & a(20) &= 0.1, & a(k) &= 0 \text{ otherwise,} \\ b(4) &= 0.2, & b(5) &= 0.1, & b(6) &= 0.7, & b(k) &= 0 \text{ otherwise.} \\ E[A] &= 7.6, & E[B] &= 5.5, & \rho &= 0.724, & c_A &= 0.598, & c_B &= 0.147. \end{aligned}$$



Hands-On: ex1/ex1.R

- *Define simple trajectory*
- *Define generators*
- *Define resources*

Can we obtain similar results as Tobias?

- ▶ **add_generator()** defines a source of events or arrivals
 - **name** is the canonical name of the resource
 - **trajectory** defines what generated events do
 - **distribution** dictates when new events occur
- ▶ **add_resource()** defines a claimable resource
 - **name** is the canonical name of the resource
 - **capacity** is the number of available processing
 - **queue_size** is the maximum number of elements allowed to wait before rejects occur
 - **mon** defines whether the resource needs to be monitored
- ▶ **trajectory()** defines a sequence of actions triggered by an event
 - Is generally processed step by step
 - Can be **rolled back** to create loops
 - Can **branch** or be **cloned**
- ▶ Basic **actions** to use in a trajectory are
 - **seize()** requests capacity from a resource
 - **timeout()** lets time pass
 - **release()** free previously requested capacity
 - **log_()** prints a log message
 - **set_attribute(), get_attribute()** sets/gets an arrival level attribute
 - **set_global(), get_global()** sets/gets a global attribute
 - **rollback()** rolls back the current trajectory by a specific amount of actions



Hands-On: ex1/ex1.R

- *Define simple trajectory*
- *Define generators*
- *Define resources*

- ▶ **add_generator()** defines a source of events or arrivals
 - **name** is the canonical name of the resource
 - **trajectory** defines what generated events do
 - **distribution** dictates when new events occur
- ▶ **add_resource()** defines a claimable resource
 - **name** is the canonical name of the resource
 - **capacity** is the number of available processing
 - **queue_size** is the maximum number of elements allowed to wait
 - **mon** defines what is monitored
- ▶ **trajectory()** defines a sequence of actions triggered by an event
 - Is generally processed step by step
 - Can be **rolled back** to create loops
 - Can **branch** or be **cloned**
- ▶ Basic **actions** to use in a trajectory are
 - **seize()** requests capacity from a resource
 - **timeout()** lets time pass

How to adapt the simulation to be a Gi/Gi/1-L with $L = 10$?
How to adapt the simulation to be an M/M/1-L model ?
What happens if $E[A] < E[B]$?



Hands-On: ex1/ex1.R

- *Define simple trajectory*
- *Define generators*
- *Define resources*



Hands-On: **video_streaming/001_basic.R**

- *Basic structure of a more complex simulation*
- *Video streaming application with limited bandwidth*
- *Clients request video based on randomized video bitrate*

Hands-On: **video_streaming/002_infrastructure.R**

- *Extension of 001_basic.R*
- *More detailed application infrastructure*
- *Frontend, backend and streaming components simulated separately*

Hands-On: **video_streaming/003_clients.R**

- *Extension of 002_infrastructure.R*
- *More detailed client behavior*
- *Clients browse, renege and exhibit rudimentary DASH behavior*